# Maestro - Python Task Orchestrator

Maestro is a production-ready Python task orchestrator that runs DAGs (Directed Acyclic Graphs) with a modern client-server architecture similar to Docker.

Features

�� Client-Server Architecture

- **REST API Server**: Persistent backend service for DAG execution
- **Lightweight CLI Client**: Docker-like command interface
- **Independent Execution**: DAGs run continuously even when client disconnects
- **Multiple Client Support**: Connect from multiple terminals to the same server

�� Core Capabilities

- **DAG-based task execution:** Define complex workflows with dependencies
- **Multi-threaded execution:** Run DAGs asynchronously with concurrent task execution
- **Multi-executor support:** Execute tasks on different platforms (local, SSH, Docker, Kubernetes)
- **Database persistence:** Track execution history and status with SQLite database
- **Real-time monitoring:** Monitor DAG execution progress with rich UI
- **Enhanced CLI interface:** Comprehensive command-line interface with monitoring and management
- **Execution lifecycle management:** Start, monitor, cancel, and cleanup DAG executions
- **YAML configuration:** Easily define your DAGs in a human-readable format
- **Extensible:** Create your own custom task types and executors
- **Comprehensive logging:** Detailed execution logs with filtering and real-time streaming
- **Status tracking:** Persistent task and DAG status with resume capabilities

�� Live Streaming

- **Real-time log streaming:** Server-Sent Events (SSE) for live log updates
- **Efficient filtering:** Client-side log deduplication and filtering
- **Graceful reconnection:** Automatic retry logic for network issues

Installation

```
# Clone the repository
git clone https://github.com/your-username/maestro.git
cd maestro

# Install dependencies (including test dependencies) using uv
uv sync --extra test

# Install the project in editable mode for local development
uv pip install -e .
```

Usage

Maestro provides two interfaces: a modern **client-server architecture** (recommended) and a legacy **standalone CLI**.

�� Client-Server Architecture (Recommended)

Starting the Server

```
# Start server in daemon mode
maestro server start --daemon --port 8000

# Start server in foreground
maestro server start --port 8000

# Check server status
maestro server status
```

Using the Client

```
# Submit a DAG for execution
maestro submit examples/sample_dag.yaml

# Check DAG status
maestro status sample_dag

# View logs
maestro logs sample_dag --limit 50

# Attach to live logs (real-time streaming)
maestro attach sample_dag

# List running DAGs
maestro running

# List all DAGs with optional filtering
maestro list --active

# Cancel a DAG
maestro cancel sample_dag

# Validate a DAG
maestro validate examples/sample_dag.yaml

# Clean up old executions
maestro cleanup --days 30
```

Server Options

```
# Custom host and port
maestro server start --host 0.0.0.0 --port 8080

# Different log levels
maestro server start --log-level debug

# Connect to remote server
maestro status sample_dag --server http://remote-host:8080
```

�� Legacy Standalone CLI

Running a DAG

```
# Run a DAG asynchronously (default)
maestro run examples/sample_dag.yaml

# Run a DAG in the background
maestro run-async examples/sample_dag.yaml
```

Validating a DAG

```
maestro validate examples/sample_dag.yaml
```

Visualizing a DAG

```
maestro visualize examples/sample_dag.yaml
```

Monitoring and Management

Monitor DAG execution in real-time

```
maestro monitor my_dag_id
```

Check DAG status

```
# Show all running DAGs
maestro status

# Show specific DAG status
maestro status my_dag_id
```

View execution logs

```
maestro logs my_dag_id
```

List all DAGs

```
# List all DAGs
maestro list-dags

# Filter by status
maestro list-dags --status running
```

View execution history

```
maestro history my_dag_id
```

Get summary statistics

```
maestro summary
```

Cancel running DAG

```
maestro cancel my_dag_id
```

Cleanup old records

```
# Clean up records older than 30 days
maestro cleanup --days 30

# Dry run to see what would be deleted
maestro cleanup --days 30 --dry-run
```

Attach to live logs

```
# Attach to live log stream
maestro attach my_dag_id

# Filter by task
maestro attach my_dag_id --task task_name

# Filter by log level
maestro attach my_dag_id --level ERROR
```

Multi-Executor Support

Maestro supports running tasks on different execution environments:

Available Executors

- **Local Executor** (default): Runs tasks on the local machine
- **SSH Executor**: Execute tasks on remote machines via SSH
- **Docker Executor**: Run tasks in Docker containers
- **Kubernetes Executor**: Execute tasks as Kubernetes jobs

Configuring Executors

Specify the executor in your DAG configuration:

```
dag:
  name: "multi_executor_dag"
  tasks:
    - task_id: "local_task"
      type: "PrintTask"
      executor: "local"  # Default
      params:
        message: "Running locally"

    - task_id: "ssh_task"
      type: "PrintTask"
      executor: "ssh"
      params:
        message: "Running via SSH"

    - task_id: "docker_task"
      type: "PrintTask"
      executor: "docker"
      params:
        message: "Running in Docker"
```

Creating Custom Executors

To create a custom executor:

1. Create a class that inherits from `maestro.core.executors.base.BaseExecutor`
2. Implement the `execute(self, task)` method
3. Register it with the executor factory

```
from maestro.core.executors.base import BaseExecutor

class CustomExecutor(BaseExecutor):
    def execute(self, task):
        # Your custom execution logic here
        task.execute_local()

# Register the executor
orchestrator.executor_factory.register_executor("custom", CustomExecutor)
```

Configuration

DAGs are defined in YAML files. Here's an example:

```
dag:
  name: "esempio_dag"
  tasks:
    - task_id: "task_1"
      type: "PrintTask"
      params:
        message: "Inizio pipeline"
        delay: 2
      dependencies: []

    - task_id: "task_2"
      type: "FileWriterTask"
      params:
        filepath: "output.txt"
```

```
      content: "Risultato elaborazione"
    dependencies: ["task_1"]

  - task_id: "task_3"
    type: "PrintTask"
    params:
      message: "Fine pipeline"
    dependencies: ["task_2"]
```

Creating Custom Tasks

To create a custom task, you need to:

1. Create a new class that inherits from `maestro.server.tasks.base.BaseTask`.
2. Define the parameters for your task as Pydantic fields.
3. Implement the execute method.
4. Register your new task in the `Orchestrator`'s `task_types` dictionary.

Architecture

Client-Server Architecture Diagram

```
sequenceDiagram
    participant CLI as CLI Client
    participant API as REST API Server
    participant DB as SQLite Database
    participant EXEC as Task Executor

    CLI->>API: POST /dags/submit
    API->>DB: Store execution metadata
    API->>EXEC: Start DAG execution (async)
    API-->>CLI: Return execution ID

    CLI->>API: GET /dags/{id}/logs/stream
    API->>DB: Query logs periodically
    API-->>CLI: Stream logs (SSE)

    CLI->>API: GET /dags/{id}/status
    API->>DB: Query execution status
    API-->>CLI: Return status

    EXEC->>DB: Update task status
    EXEC->>DB: Write execution logs
```

Interface Comparison

Feature	Client-Server	Legacy CLI
Execution Model	Persistent server	Process-based
DAG Persistence	Survives client disconnect	Dies with process
Multi-client Support	Multiple clients	Single process
Live Log Streaming	Real-time SSE	Terminal attach
Remote Access	Network accessible	Local only
Resource Usage	Lower client overhead	Higher memory usage
Deployment	Production ready	Development/testing

Class Diagram

```
classDiagram
    direction LR
    class Task {
        <<abstract>>
        +str task_id
        +List~str~ dependencies
        +TaskStatus status
        +Callable on_success
        +Callable on_failure
        +execute()
    }
```

```
class BaseTask {
    +execute()
}

class PrintTask {
    +str message
    +Optional~int~ delay
    +execute()
}

class FileWriterTask {
    +str filepath
    +str content
    +Literal~"append", "overwrite"~ mode
    +execute()
}

class WaitTask {
    +int delay
    +execute()
}

class DAG {
    +Dict~str, Task~ tasks
    +add_task(task: Task)
    +validate()
    +get_execution_order() List~str~
    +execute()
}

class Orchestrator {
    +Dict~str, type~ task_types
    +load_dag_from_file(filepath: str) DAG
    +run_dag(dag: DAG)
    +visualize_dag(dag: DAG)
    +get_dag_status(dag: DAG) Dict~str, Any~
}

class APIServer {
    +FastAPI app
    +Orchestrator orchestrator
    +submit_dag(request)
    +get_dag_status(dag_id)
    +stream_logs(dag_id)
    +get_running_dags()
}

class APIClient {
    +str base_url
    +submit_dag(dag_file_path)
    +get_dag_status(dag_id)
    +stream_dag_logs(dag_id)
    +get_running_dags()
}

Task <|-- BaseTask
BaseTask <|-- PrintTask
BaseTask <|-- FileWriterTask
BaseTask <|-- WaitTask
DAG "1" -- "*" Task : contains
Orchestrator ..> DAGLoader : uses
DAGLoader ..> DAG : creates
Orchestrator ..> Task : manages
APIServer ..> Orchestrator : uses
APIClient ..> APIServer : HTTP calls
```

Testing

Maestro includes a comprehensive test suite covering all major functionality:

Test Categories

Core Functionality Tests

- **DAG Operations**: DAG creation, validation, cycle detection
- **Task Execution**: Task lifecycle, dependencies, status tracking
- **Database Features**: Persistence, state management, resume functionality

Multi-Executor Tests

- **Executor Factory**: Registration, retrieval, thread safety
- **Custom Executors**: Creation, registration, isolation
- **Executor Integration**: Orchestrator integration, error handling

Enhanced CLI Tests

- **Status Management**: DAG execution tracking, history, summaries
- **Database Operations**: Cleanup, cancellation, log management
- **CLI Integration**: Command scenarios, data format validation

Performance & Concurrency Tests

- **Thread Safety**: Concurrent status updates, database access
- **Async Execution**: Background DAG execution, monitoring
- **Resource Management**: Database connection handling, cleanup

Running Tests

```
# Run all tests
./run_tests.sh

# Run specific test categories
uv run pytest tests/test_multi_executor.py -v       # Multi-executor tests
uv run pytest tests/test_enhanced_cli.py -v        # Enhanced CLI tests
uv run pytest tests/test_db_feature.py -v          # Database features
uv run pytest tests/test_dag.py -v                 # DAG operations
uv run pytest tests/test_orchestrator_dagloader.py -v # Orchestrator tests

# Run tests with coverage
uv run pytest --cov=maestro --cov-report=html
```

Test Coverage

- **Multi-executor support**: 13/13 tests passing
- **Enhanced CLI features**: 15/15 tests passing (non-concurrent)
- **Database operations**: 4/4 tests passing
- **DAG functionality**: 3/3 tests passing
- **Orchestrator features**: 5/5 tests passing
- **Concurrent execution**: Limited by SQLite threading constraints

Known Limitations

- **SQLite Threading**: Concurrent database access may cause issues in heavy multi-threading scenarios
- **Test Isolation**: Some concurrent tests may experience timing-related failures

Continuous Integration

For production use, consider: - Using PostgreSQL or MySQL for better concurrent access - Implementing connection pooling for database operations - Adding integration tests for specific deployment environments