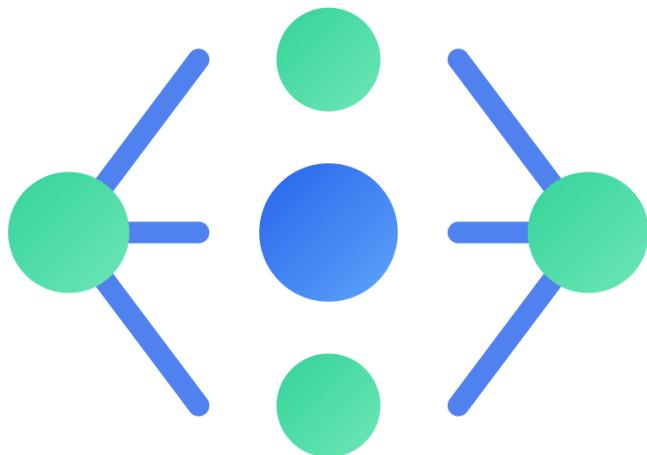


# Maestro - Python Task Orchestrator



License [MIT](#) Python [3.13](#)

Maestro is a production-ready Python task orchestrator that runs DAGs (Directed Acyclic Graphs) with a modern client-server architecture similar to Docker.

## Features

### Client-Server Architecture

- **REST API Server:** Persistent backend service for DAG execution
- **Lightweight CLI Client:** Docker-like command interface
- **Independent Execution:** DAGs run continuously even when client disconnects
- **Multiple Client Support:** Connect from multiple terminals to the same server

### Core Capabilities

- **DAG-based task execution:** Define complex workflows with dependencies
- **Multi-threaded execution:** Run DAGs asynchronously with concurrent task execution
- **Multi-executor support:** Execute tasks on different platforms (local, SSH, Docker, Kubernetes)
- **Database persistence:** Track execution history and status with SQLite database
- **Real-time monitoring:** Monitor DAG execution progress with rich UI
- **Enhanced CLI interface:** Comprehensive command-line interface with monitoring and management
- **Execution lifecycle management:** Start, monitor, cancel, and cleanup DAG executions
- **YAML configuration:** Easily define your DAGs in a human-readable format
- **Extensible:** Create your own custom task types and executors
- **Comprehensive logging:** Detailed execution logs with filtering and real-time streaming
- **Status tracking:** Persistent task and DAG status with resume capabilities

### Live Streaming

- **Real-time log streaming:** Server-Sent Events (SSE) for live log updates
- **Efficient filtering:** Client-side log deduplication and filtering
- **Graceful reconnection:** Automatic retry logic for network issues

## Installation and use of UV

```
# Clone the repository
git clone https://github.com/your-username/maestro.git
cd maestro

# Install dependencies (including test dependencies) using uv
uv sync --extra test

# Install the project in editable mode for local development
uv pip install -e .

# ...

# NOTE: Since uv provides and manages all dependencies required by the
# software, it's recommended to use uv run when executing commands such as
# tests or CLI operations.
# This ensures that all dependencies are correctly resolved within the
# managed environment.

# Example:
uv run maestro server start
uv run pytest

# Alternatively, you can start an interactive environment once and run
# commands normally (if "uv shell" is implemented):
uv shell
maestro server start
pytest
```

## Usage

Maestro provides two interfaces: a modern **client-server architecture** (recommended) and a legacy **standalone CLI**.

### Client-Server Architecture (Recommended)

#### Starting the Server

```
# Start server in daemon mode
maestro server start --daemon --port 8000

# Start server in foreground
maestro server start --port 8000

# Check server status
maestro server status
```

## Using the Client

```
# Submit a DAG for execution
maestro submit examples/sample_dag.yaml

# Check DAG status
maestro status sample_dag

# View logs
maestro logs sample_dag --limit 50

# Attach to live logs (real-time streaming)
maestro attach sample_dag

# List running DAGs
maestro running

# List all DAGs with optional filtering
maestro list --active

# Cancel a DAG
maestro cancel sample_dag

# Validate a DAG
maestro validate examples/sample_dag.yaml

# Clean up old executions
maestro cleanup --days 30
```

## Server Options

```
# Custom host and port
maestro server start --host 0.0.0.0 --port 8080

# Different log levels
maestro server start --log-level debug

# Connect to remote server
maestro status sample_dag --server http://remote-host:8080
```

## Legacy Standalone CLI

## Running a DAG

```
# Run a DAG asynchronously (default)
maestro run examples/sample_dag.yaml
```

```
# Run a DAG in the background  
maestro run-async examples/sample_dag.yaml
```

## Validating a DAG

```
maestro validate examples/sample_dag.yaml
```

## Visualizing a DAG

```
maestro visualize examples/sample_dag.yaml
```

## Monitoring and Management

### Monitor DAG execution in real-time

```
maestro monitor my_dag_id
```

### Check DAG status

```
# Show all running DAGs  
maestro status  
  
# Show specific DAG status  
maestro status my_dag_id
```

### View execution logs

```
maestro logs my_dag_id
```

### List all DAGs

```
# List all DAGs  
maestro list-dags  
  
# Filter by status  
maestro list-dags --status running
```

## View execution history

```
maestro history my_dag_id
```

## Get summary statistics

```
maestro summary
```

## Cancel running DAG

```
maestro cancel my_dag_id
```

## Cleanup old records

```
# Clean up records older than 30 days
maestro cleanup --days 30

# Dry run to see what would be deleted
maestro cleanup --days 30 --dry-run
```

## Attach to live logs

```
# Attach to live log stream
maestro attach my_dag_id

# Filter by task
maestro attach my_dag_id --task task_name

# Filter by log level
maestro attach my_dag_id --level ERROR
```

## Multi-Executor Support

Maestro supports running tasks on different execution environments:

### Available Executors

- **Local Executor** (default): Runs tasks on the local machine
- **SSH Executor**: Execute tasks on remote machines via SSH
- **Docker Executor**: Run tasks in Docker containers

- **Kubernetes Executor:** Execute tasks as Kubernetes jobs

## Configuring Executors

Specify the executor in your DAG configuration:

```
dag:  
  name: "multi_executor_dag"  
  tasks:  
    - task_id: "local_task"  
      type: "PrintTask"  
      executor: "local" # Default  
      params:  
        message: "Running locally"  
  
    - task_id: "ssh_task"  
      type: "PrintTask"  
      executor: "ssh"  
      params:  
        message: "Running via SSH"  
  
    - task_id: "docker_task"  
      type: "PrintTask"  
      executor: "docker"  
      params:  
        message: "Running in Docker"
```

## Creating Custom Executors

To create a custom executor:

1. Create a class that inherits from `maestro.core.executors.base.BaseExecutor`
2. Implement the `execute(self, task)` method
3. Register it with the executor factory

```
from maestro.core.executors.base import BaseExecutor  
  
class CustomExecutor(BaseExecutor):  
    def execute(self, task):  
        # Your custom execution logic here  
        task.execute_local()  
  
    # Register the executor  
    orchestrator.executor_factory.register_executor("custom", CustomExecutor)
```

## Configuration

DAGs are defined in YAML files. Here's an example:

```

dag:
  name: "esempio_dag"
  tasks:
    - task_id: "task_1"
      type: "PrintTask"
      params:
        message: "Inizio pipeline"
        delay: 2
      dependencies: []

    - task_id: "task_2"
      type: "FileWriterTask"
      params:
        filepath: "output.txt"
        content: "Risultato elaborazione"
      dependencies: ["task_1"]

    - task_id: "task_3"
      type: "PrintTask"
      params:
        message: "Fine pipeline"
      dependencies: ["task_2"]

```

## Creating Custom Tasks

To create a custom task, you need to:

1. Create a new class that inherits from `maestro.server.tasks.base.BaseTask`.
2. Define the parameters for your task as Pydantic fields.
3. Implement the `execute` method.
4. Register your new task in the `Orchestrator`'s `task_types` dictionary.

## Architecture

### Client-Server Architecture Diagram

```

sequenceDiagram
    participant CLI as CLI Client
    participant API as REST API Server
    participant DB as SQLite Database
    participant EXEC as Task Executor

    CLI->>API: POST /dags/submit
    API->>DB: Store execution metadata
    API->>EXEC: Start DAG execution (async)
    API-->>CLI: Return execution ID

    CLI->>API: GET /dags/{id}/logs/stream
    API->>DB: Query logs periodically
    API-->>CLI: Stream logs (SSE)

```

```
CLI->>API: GET /dags/{id}/status
API->>DB: Query execution status
API-->>CLI: Return status
```

```
EXEC->>DB: Update task status
EXEC->>DB: Write execution logs
```

## Interface Comparison

Feature	Client-Server	Legacy CLI
<b>Execution Model</b>	Persistent server	Process-based
<b>DAG Persistence</b>	✓ Survives client disconnect	✗ Dies with process
<b>Multi-client Support</b>	✓ Multiple clients	✗ Single process
<b>Live Log Streaming</b>	✓ Real-time SSE	✓ Terminal attach
<b>Remote Access</b>	✓ Network accessible	✗ Local only
<b>Resource Usage</b>	Lower client overhead	Higher memory usage
<b>Deployment</b>	Production ready	Development/testing

## Class Diagram

```
classDiagram
    direction LR
    class Task {
        <>
        +str task_id
        +List~str~ dependencies
        +TaskStatus status
        +Callable on_success
        +Callable on_failure
        +execute()
    }

    class BaseTask {
        +execute()
    }

    class PrintTask {
        +str message
        +Optional~int~ delay
        +execute()
    }

    class FileWriterTask {
        +str filepath
        +str content
    }
```

```
+Literal~"append", "overwrite"~ mode
+execute()
}

class WaitTask {
    +int delay
    +execute()
}

class DAG {
    +Dict~str, Task~ tasks
    +add_task(task: Task)
    +validate()
    +get_execution_order() List~str~
    +execute()
}

class Orchestrator {
    +Dict~str, type~ task_types
    +load_dag_from_file(filepath: str) DAG
    +run_dag(dag: DAG)
    +visualize_dag(dag: DAG)
    +get_dag_status(dag: DAG) Dict~str, Any~
}

class APIServer {
    +FastAPI app
    +Orchestrator orchestrator
    +submit_dag(request)
    +get_dag_status(dag_id)
    +stream_logs(dag_id)
    +get_running_dags()
}

class APIClient {
    +str base_url
    +submit_dag(dag_file_path)
    +get_dag_status(dag_id)
    +stream_dag_logs(dag_id)
    +get_running_dags()
}

Task <|-- BaseTask
BaseTask <|-- PrintTask
BaseTask <|-- FileWriterTask
BaseTask <|-- WaitTask
DAG "1" -- "*" Task : contains
Orchestrator ..> DAGLoader : uses
DAGLoader ..> DAG : creates
Orchestrator ..> Task : manages
APIServer ..> Orchestrator : uses
APIClient ..> APIServer : HTTP calls
```

# Testing

Maestro includes a comprehensive test suite covering all major functionality:

## Test Categories

### Core Functionality Tests

- **DAG Operations:** DAG creation, validation, cycle detection
- **Task Execution:** Task lifecycle, dependencies, status tracking
- **Database Features:** Persistence, state management, resume functionality

### Multi-Executor Tests

- **Executor Factory:** Registration, retrieval, thread safety
- **Custom Executors:** Creation, registration, isolation
- **Executor Integration:** Orchestrator integration, error handling

### Enhanced CLI Tests

- **Status Management:** DAG execution tracking, history, summaries
- **Database Operations:** Cleanup, cancellation, log management
- **CLI Integration:** Command scenarios, data format validation

### Performance & Concurrency Tests

- **Thread Safety:** Concurrent status updates, database access
- **Async Execution:** Background DAG execution, monitoring
- **Resource Management:** Database connection handling, cleanup

## Running Tests

```
# Run all tests
./run_tests.sh

# Run specific test categories
uv run pytest tests/test_multi_executor.py -v          # Multi-executor
tests
uv run pytest tests/test_enhanced_cli.py -v           # Enhanced CLI tests
uv run pytest tests/test_db_feature.py -v             # Database features
uv run pytest tests/test_dag.py -v                    # DAG operations
uv run pytest tests/test_orchestrator_dagloader.py -v # Orchestrator tests

# Run tests with coverage
uv run pytest --cov=maestro --cov-report=html
```

## Test Coverage

- **Multi-executor support:** 13/13 tests passing

- **✓ Enhanced CLI features:** 15/15 tests passing (non-concurrent)
- **✓ Database operations:** 4/4 tests passing
- **✓ DAG functionality:** 3/3 tests passing
- **✓ Orchestrator features:** 5/5 tests passing
- **⚠ Concurrent execution:** Limited by SQLite threading constraints

## Known Limitations

- **SQLite Threading:** Concurrent database access may cause issues in heavy multi-threading scenarios
- **Test Isolation:** Some concurrent tests may experience timing-related failures

## Continuous Integration

For production use, consider:

- Using PostgreSQL or MySQL for better concurrent access
- Implementing connection pooling for database operations
- Adding integration tests for specific deployment environments